

Multiple Regression Project

Kabin Devkota, Nishan Khanal and Udit Bista

4/15/2025

Introduction

This analysis focuses on a data set from Kaggle that describes the nutrition facts for McDonald's Menu. This dataset provides a nutrition analysis of every menu item on the US McDonald's menu, including breakfast, beef burgers, chicken and fish sandwiches, fries, salads, soda, coffee and tea, milkshakes, and desserts. The data for this analysis consists of response variable $y = \text{Calories}$ and following explanatory variables:

Category

item

ServingSize

TotalFat

TotalFat(%DailyValue)

SaturatedFat

SaturatedFat(%DailyValue)

TransFat

Cholesterol

Cholesterol(%DailyValue)

Sodium

Sodium(%DailyValue)

Carbohydrates

Carbohydrates(%DailyValue)

DietaryFiber

DietaryFiber(%DailyValue)

Sugars

Protein

VitaminA(%DailyValue)

VitaminC(%DailyValue)

Calcium(%DailyValue)

Iron(%DailyValue)

```
# library(ggplot2)
library(lmtest, pos=4)
```

```
## Loading required package: zoo

##
## Attaching package: 'zoo'

## The following objects are masked from 'package:base':
##
##      as.Date, as.Date.numeric

library(corrplot)

## corrplot 0.95 loaded

# reading the data from the csv file which has some values containing comma enclosed by double quotes
mcdonalds = read.csv("Mcdonalds_menu.csv",header=TRUE,quote="\"",sep=",")
head(mcdonalds)
```

##	Category	Item	Serving.Size	Calories
## 1	Breakfast	Egg McMuffin	4.8 oz (136 g)	300
## 2	Breakfast	Egg White Delight	4.8 oz (135 g)	250
## 3	Breakfast	Sausage McMuffin	3.9 oz (111 g)	370
## 4	Breakfast	Sausage McMuffin with Egg	5.7 oz (161 g)	450
## 5	Breakfast	Sausage McMuffin with Egg Whites	5.7 oz (161 g)	400
## 6	Breakfast	Steak & Egg McMuffin	6.5 oz (185 g)	430

##	Calories.from.Fat	Total.Fat	Total.Fat....Daily.Value.	Saturated.Fat
## 1	120	13	20	5
## 2	70	8	12	3
## 3	200	23	35	8
## 4	250	28	43	10
## 5	210	23	35	8
## 6	210	23	36	9

##	Saturated.Fat....Daily.Value.	Trans.Fat	Cholesterol
## 1	25	0	260
## 2	15	0	25
## 3	42	0	45
## 4	52	0	285
## 5	42	0	50
## 6	46	1	300

##	Cholesterol....Daily.Value.	Sodium	Sodium....Daily.Value.	Carbohydrates
## 1	87	750	31	31
## 2	8	770	32	30
## 3	15	780	33	29
## 4	95	860	36	30
## 5	16	880	37	30
## 6	100	960	40	31

##	Carbohydrates....Daily.Value.	Dietary.Fiber	Dietary.Fiber....Daily.Value.
## 1	10	4	17
## 2	10	4	17
## 3	10	4	17
## 4	10	4	17
## 5	10	4	17
## 6	10	4	18

##	Sugars	Protein	Vitamin.A....Daily.Value.	Vitamin.C....Daily.Value.
## 1	3	17	10	0
## 2	3	18	6	0
## 3	2	14	8	0
## 4	2	21	15	0

```
## 5      2      21      6      0
## 6      3      26     15      2
##  Calcium....Daily.Value.  Iron....Daily.Value.
## 1              25              15
## 2              25              8
## 3              25             10
## 4              30             15
## 5              25             10
## 6              30             20
```

```
# getting the column names
colnames(mcdonalds)
```

```
## [1] "Category"           "Item"
## [3] "Serving.Size"       "Calories"
## [5] "Calories.from.Fat"  "Total.Fat"
## [7] "Total.Fat....Daily.Value." "Saturated.Fat"
## [9] "Saturated.Fat....Daily.Value." "Trans.Fat"
## [11] "Cholesterol"        "Cholesterol....Daily.Value."
## [13] "Sodium"             "Sodium....Daily.Value."
## [15] "Carbohydrates"      "Carbohydrates....Daily.Value."
## [17] "Dietary.Fiber"      "Dietary.Fiber....Daily.Value."
## [19] "Sugars"             "Protein"
## [21] "Vitamin.A....Daily.Value." "Vitamin.C....Daily.Value."
## [23] "Calcium....Daily.Value." "Iron....Daily.Value."
```

```
# cleaning column names: replace spaces, %, parentheses, etc.
colnames(mcdonalds) <- make.names(colnames(mcdonalds))
```

Out of these 22 explanatory variables, some of the variables like “category”, “Total.Fat....Daily.Value.”, “Cholesterol....Daily.Value.”, etc are either redundant because there already exist other columns that account for those values or irrelevant for prediction of calories in food items. Therefore, we are dropping them from the analysis to avoid multicollinearity and improve model efficiency.

```
# dropping the columns that is irrelevant or redundant for the analysis
cols_to_drop <- c(
  "Category",
  "Calories.from.Fat",
  "Saturated.Fat",
  "Trans.Fat",
  "Total.Fat....Daily.Value.",
  "Saturated.Fat....Daily.Value.",
  "Cholesterol....Daily.Value.",
  "Sodium....Daily.Value.",
  "Carbohydrates....Daily.Value.",
  "Sugars",
  "Dietary.Fiber....Daily.Value."
)
mcdonalds <- mcdonalds[, !(names(mcdonalds) %in% cols_to_drop)]
```

Now, we are converting the % Daily Value columns to absolute values. The daily values are based on a 2000 calorie diet, and the conversion is done using the following formula:

$$\text{Absolute Value} = \left(\frac{\text{Percentage Daily Value}}{100} \right) \times \text{Daily Value}$$

where the daily values are as follows:

```

# Daily values (units must match dataset)
daily_values <- c(
  "Vitamin.A" = 900,    # mcg RAE
  "Vitamin.C" = 90,     # mg
  "Calcium" = 1300,     # mg
  "Iron" = 18           # mg
)

# Map of columns to their associated nutrients
conversion_map <- list(
  "Vitamin.A....Daily.Value." = "Vitamin.A",
  "Vitamin.C....Daily.Value." = "Vitamin.C",
  "Calcium....Daily.Value." = "Calcium",
  "Iron....Daily.Value." = "Iron"
)

# For each %DV column, calculate the absolute value and overwrite the same column with absolute value
for (dv_col in names(conversion_map)) {
  nutrient <- conversion_map[[dv_col]]

  if (dv_col %in% names(mcdonalds)) {
    # Create a new column name or overwrite the existing one
    new_col <- nutrient # Replace the DV column with nutrient name only
    mcdonalds[[new_col]] <- (mcdonalds[[dv_col]] / 100) * daily_values[[nutrient]]
  }
}

```

The values in the serving size column contain both numbers and units (e.g., “15 oz”). In this step, we are extracting only the numeric part and discarding the unit, so “15 oz” becomes just 15.

```

# Extract numeric part before "oz" or "fl oz" from the Serving Size column
mcdonalds$Serving.Size.Oz <- as.numeric(sub("([0-9.]+)\\s*(fl\\s*)?oz.*", "\\1", mcdonalds$Serving.Size))

```

Warning: NAs introduced by coercion

```
head(mcdonalds)
```

```

##           Item Serving.Size Calories Total.Fat
## 1      Egg McMuffin 4.8 oz (136 g)      300      13
## 2      Egg White Delight 4.8 oz (135 g)      250      8
## 3      Sausage McMuffin 3.9 oz (111 g)      370     23
## 4 Sausage McMuffin with Egg 5.7 oz (161 g)      450     28
## 5 Sausage McMuffin with Egg Whites 5.7 oz (161 g)      400     23
## 6      Steak & Egg McMuffin 6.5 oz (185 g)      430     23
## Cholesterol Sodium Carbohydrates Dietary.Fiber Protein
## 1      260      750           31           4      17
## 2       25      770           30           4      18
## 3       45      780           29           4      14
## 4      285      860           30           4      21
## 5       50      880           30           4      21
## 6      300      960           31           4      26
## Vitamin.A....Daily.Value. Vitamin.C....Daily.Value. Calcium....Daily.Value.
## 1              10              0              25
## 2               6              0              25
## 3               8              0              25

```

```
## 4          15          0          30
## 5          6          0          25
## 6          15          2          30
##   Iron....Daily.Value. Vitamin.A Vitamin.C Calcium Iron Serving.Size.Oz
## 1          15          90          0.0      325 2.70          4.8
## 2           8          54          0.0      325 1.44          4.8
## 3          10          72          0.0      325 1.80          3.9
## 4          15         135          0.0      390 2.70          5.7
## 5          10          54          0.0      325 1.80          5.7
## 6          20         135          1.8      390 3.60          6.5
```

```
# dropping the data points with null values that would be a hindrance for the analysis
mcdonalds <- na.omit(mcdonalds)
```

```
head(mcdonalds)
```

```
##           Item   Serving.Size Calories Total.Fat
## 1      Egg McMuffin 4.8 oz (136 g)      300      13
## 2    Egg White Delight 4.8 oz (135 g)      250       8
## 3      Sausage McMuffin 3.9 oz (111 g)      370      23
## 4 Sausage McMuffin with Egg 5.7 oz (161 g)      450      28
## 5 Sausage McMuffin with Egg Whites 5.7 oz (161 g)      400      23
## 6    Steak & Egg McMuffin 6.5 oz (185 g)      430      23
## Cholesterol Sodium Carbohydrates Dietary.Fiber Protein
## 1          260      750          31          4      17
## 2           25      770          30          4      18
## 3           45      780          29          4      14
## 4          285      860          30          4      21
## 5           50      880          30          4      21
## 6          300      960          31          4      26
##   Vitamin.A....Daily.Value. Vitamin.C....Daily.Value. Calcium....Daily.Value.
## 1          10          0          25
## 2           6          0          25
## 3           8          0          25
## 4          15          0          30
## 5           6          0          25
## 6          15          2          30
##   Iron....Daily.Value. Vitamin.A Vitamin.C Calcium Iron Serving.Size.Oz
## 1          15          90          0.0      325 2.70          4.8
## 2           8          54          0.0      325 1.44          4.8
## 3          10          72          0.0      325 1.80          3.9
## 4          15         135          0.0      390 2.70          5.7
## 5          10          54          0.0      325 1.80          5.7
## 6          20         135          1.8      390 3.60          6.5
```

```
# Identify all columns to normalize (exclude Item and Serving Size Oz)
```

```
cols_to_normalize <- setdiff(names(mcdonalds), c("Item", "Serving.Size", "Serving.Size.Oz"))
```

```
# Convert values to per oz
```

```
mcdonalds[cols_to_normalize] <- lapply(mcdonalds[cols_to_normalize], function(col) col / mcdonalds$Serving.Size.Oz)
```

```
# dropping the remaining columns containing daily value percentages and serving size
```

```
cols_to_drop <- c(
  "Vitamin.A....Daily.Value.",
  "Vitamin.C....Daily.Value.",
```

```

"Calcium...Daily.Value.",
"Iron...Daily.Value.",
"Serving.Size"
)
mcdonalds <- mcdonalds[, !(names(mcdonalds) %in% cols_to_drop)]

# saving all the column names except "Item" and "Serving Size Oz"
col_names <- setdiff(names(mcdonalds), c("Item", "Serving.Size.Oz"))

col_names

```

```

## [1] "Calories"      "Total.Fat"      "Cholesterol"    "Sodium"
## [5] "Carbohydrates" "Dietary.Fiber"  "Protein"        "Vitamin.A"
## [9] "Vitamin.C"     "Calcium"        "Iron"

```

After the completion of the data preprocessing part, our data set has been narrowed down to one response variable *calories* and ten explanatory variables:

1. $x_1 = TotalFat$
2. $x_2 = Cholestrol$
3. $x_3 = Sodium$
4. $x_4 = Carbohydrates$
5. $x_5 = DietaryFiber$
6. $x_6 = Protein$
7. $x_7 = VitaminA$
8. $x_8 = VitaminC$
9. $x_9 = Calcium$
10. $x_{10} = Iron$

```

# creating new names: y for Calories, x1, x2, ... for the rest
new_names <- c("y", paste0("x", seq_along(col_names[-which(col_names == "Calories")])))

# creating a named vector to rename the columns
name_map <- setNames(new_names, c("Calories", col_names[col_names != "Calories"]))

# renaming the columns
names(mcdonalds)[names(mcdonalds) %in% names(name_map)] <- name_map[names(mcdonalds)[names(mcdonalds) %
# saving all the column names except "Item" and "Serving Size Oz"
col_names <- setdiff(names(mcdonalds), c("Item", "Serving.Size.Oz"))

```

Correlation Coefficients

Since we are looking at linear relationships between the outcome variable (*calories*) with each explanatory variable, it may be of interest to determine the correlation coefficients between the outcome variable with each explanatory variable.

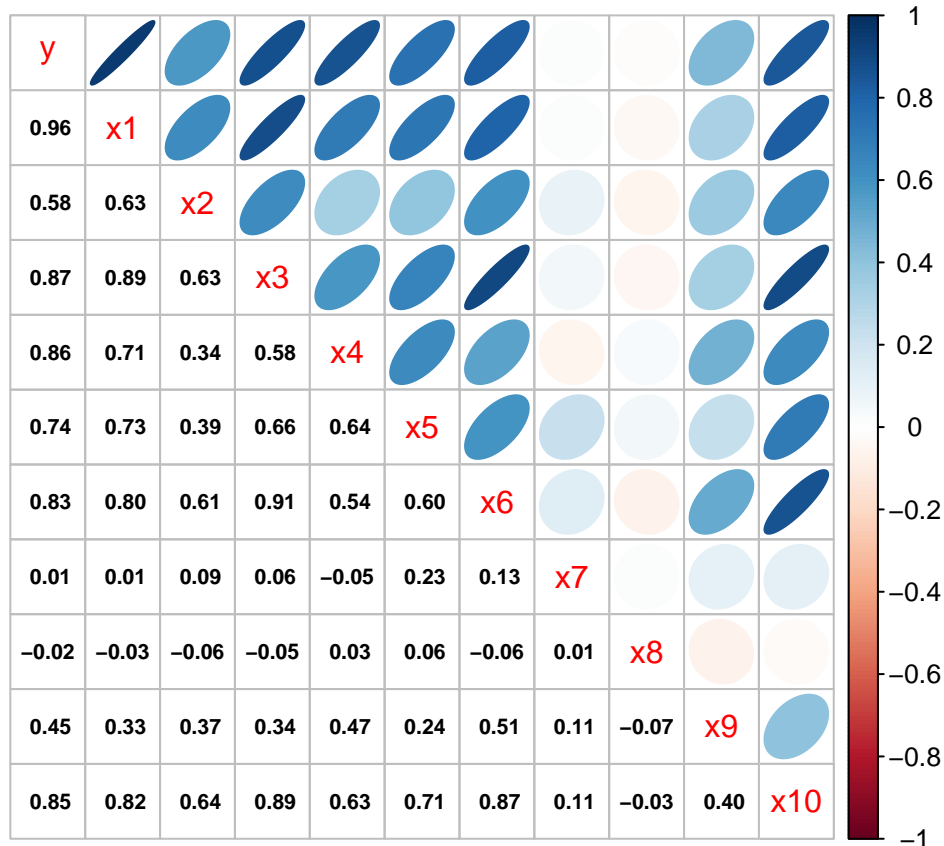
```

corr_matrix <- cor(mcdonalds[, col_names], use="everything")
round(corr_matrix, 3)

```

```
##          y      x1      x2      x3      x4      x5      x6      x7      x8      x9      x10
## y      1.000  0.960  0.579  0.872  0.862  0.745  0.826  0.012 -0.019  0.448  0.848
## x1      0.960  1.000  0.628  0.890  0.709  0.726  0.802  0.010 -0.035  0.326  0.823
## x2      0.579  0.628  1.000  0.626  0.336  0.394  0.607  0.092 -0.056  0.367  0.642
## x3      0.872  0.890  0.626  1.000  0.581  0.664  0.906  0.056 -0.047  0.339  0.890
## x4      0.862  0.709  0.336  0.581  1.000  0.638  0.538 -0.050  0.033  0.470  0.631
## x5      0.745  0.726  0.394  0.664  0.638  1.000  0.599  0.229  0.057  0.237  0.709
## x6      0.826  0.802  0.607  0.906  0.538  0.599  1.000  0.134 -0.063  0.510  0.868
## x7      0.012  0.010  0.092  0.056 -0.050  0.229  0.134  1.000  0.013  0.107  0.114
## x8     -0.019 -0.035 -0.056 -0.047  0.033  0.057 -0.063  0.013  1.000 -0.067 -0.025
## x9      0.448  0.326  0.367  0.339  0.470  0.237  0.510  0.107 -0.067  1.000  0.403
## x10     0.848  0.823  0.642  0.890  0.631  0.709  0.868  0.114 -0.025  0.403  1.000
```

```
corrplot.mixed(corr_matrix, lower.col = "black", number.cex = .7, upper = "ellipse")
```



Multiple Regression

At this point we have interest in building a model for calories using some combination of the explanatory variables. Using multiple regression one initially build a model with all of the possible explanatory variables. Below is some R output for this Multiple Linear Regression (MLR) analysis.

General Form for a Multiple Regression Model

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_3 + \cdots + \beta_{14} X_{14}$$

```
model.1 <- lm(y~x1+x2+x3+x4+x5+x6+x7+x8+x9+x10, data=mcDonalds)
```

```
summary(model.1)
```

```
##
## Call:
## lm(formula = y ~ x1 + x2 + x3 + x4 + x5 + x6 + x7 + x8 + x9 +
##      x10, data = mcdonalds)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.49593 -0.29807  0.04731  0.31380  2.39670
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.120801   0.084412  -1.431  0.15368
## x1           9.066437   0.061215 148.109 < 2e-16 ***
## x2          -0.014724   0.005112  -2.880  0.00433 **
## x3           0.002323   0.001659   1.400  0.16282
## x4           4.031077   0.024589 163.937 < 2e-16 ***
## x5          -0.719929   0.275351  -2.615  0.00949 **
## x6           3.776797   0.077650  48.639 < 2e-16 ***
## x7           0.001713   0.001739   0.985  0.32571
## x8          -0.002740   0.004894  -0.560  0.57608
## x9          -0.001219   0.003086  -0.395  0.69325
## x10          0.922310   0.455764   2.024  0.04409 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.6007 on 245 degrees of freedom
## Multiple R-squared:  0.9996, Adjusted R-squared:  0.9996
## F-statistic: 6.146e+04 on 10 and 245 DF, p-value: < 2.2e-16
```

Equation of the Model with all of the Explanatory Variables

(Note: This is referred to as the Full Model)

$$Y = -0.120 + 9.066X_1 - 0.014X_2 + 0.002X_3 + 4.031X_4 + \cdots + 0.922X_{10}$$

Coefficient of Determination

Interpretation: **99.96** % of the variability in calories is accounted for in this model, (i.e., is accounted for in the model between calories and the thirteen explanatory variables).

Test for the Significance of the Model:

Ho: None of the explanatory variables is a linear predictor of calories (i.e., the model is not significant or is not useful in predicting the response)

Ha: At least one of the explanatory variables is a significant linear predictor of calories (i.e., the model is significant or at least some portion of the model is useful in predicting the response)

Test statistic: $F^* = 6.146e+04$

P-value: $< 2.2e-16$

Conclusion: Reject Ho in favor of Ha. There is sufficient evidence to conclude that at least one of the explanatory variables is a significant linear predictor of calories (i.e., the model is significant or at least some portion of the model is useful in predicting the response)

Further Analysis

Since at least one of the independent variables is significant, we do further analysis to determine which one(s) is/are significant.

Test for an Individual Predictor in this Model:

Ho: With x1, x2, x3, x4, x5, x6, x7, x8, and x10 in the model, x9 is not a linear predictor of y

Ha: With x1, x2, x3, x4, x5, x6, x7, x8, and x10 in the model, x9 is a significant linear predictor of y

Test statistic: $t^* = -0.395$

P-value: 0.69325

Conclusion: Fail to reject Ho. There is insufficient evidence to conclude that with x1, x2, x3, x4, x5, x6, x7, x8, and x10 in the model, x9 is not a linear predictor of y

In the above model, x9 has the largest p-value and thus is the least significant. We could remove it from the model and rerun the analysis. Then we could test for significance of another independent variable. We could continue this process until only significant variables are left. This method for identifying the best model is referred to as **Backward Selection**.

Some selected output for the **Backward Selection** procedure:

```
model.2 <- lm(y~x1+x2+x3+x4+x5+x6+x7+x8+x10, data=mcdonalds)

summary(model.2)

##
## Call:
## lm(formula = y ~ x1 + x2 + x3 + x4 + x5 + x6 + x7 + x8 + x10,
##     data = mcdonalds)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.48249 -0.30917  0.04845  0.30838  2.38516
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.125508   0.083422  -1.504  0.13374
## x1           9.071362   0.059827 151.625 < 2e-16 ***
## x2          -0.015260   0.004920  -3.102  0.00215 **
## x3           0.002475   0.001611   1.536  0.12593
## x4           4.026415   0.021533 186.989 < 2e-16 ***
## x5          -0.709413   0.273590  -2.593  0.01008 *
## x6           3.761373   0.066997  56.142 < 2e-16 ***
## x7           0.001667   0.001733   0.962  0.33681
## x8          -0.002618   0.004875  -0.537  0.59181
## x10          0.936844   0.453496   2.066  0.03989 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.5997 on 246 degrees of freedom
## Multiple R-squared:  0.9996, Adjusted R-squared:  0.9996
## F-statistic: 6.852e+04 on 9 and 246 DF, p-value: < 2.2e-16
```

```
model.3 <- lm(y~x1+x2+x3+x4+x5+x6+x7+x10, data=mcdonalds)
```

```
summary(model.3)
```

```
##
## Call:
## lm(formula = y ~ x1 + x2 + x3 + x4 + x5 + x6 + x7 + x10, data = mcdonalds)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-2.47965	-0.30666	0.05207	0.30998	2.39673

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.128464	0.083120	-1.546	0.12350
x1	9.072878	0.059675	152.039	< 2e-16 ***
x2	-0.015233	0.004913	-3.101	0.00215 **
x3	0.002468	0.001609	1.534	0.12642
x4	4.025910	0.021481	187.415	< 2e-16 ***
x5	-0.722892	0.272043	-2.657	0.00839 **
x6	3.763033	0.066829	56.308	< 2e-16 ***
x7	0.001671	0.001730	0.966	0.33515
x10	0.935256	0.452832	2.065	0.03993 *

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.5988 on 247 degrees of freedom
## Multiple R-squared:  0.9996, Adjusted R-squared:  0.9996
## F-statistic: 7.731e+04 on 8 and 247 DF,  p-value: < 2.2e-16
```

```
model.4 <- lm(y~x1+x2+x3+x4+x5+x6+x10, data=mcdonalds)
```

```
summary(model.4)
```

```
##
## Call:
## lm(formula = y ~ x1 + x2 + x3 + x4 + x5 + x6 + x10, data = mcdonalds)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-2.49076	-0.29629	0.05666	0.30846	2.36800

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.109634	0.080790	-1.357	0.17601
x1	9.064833	0.059082	153.427	< 2e-16 ***
x2	-0.014810	0.004892	-3.027	0.00273 **
x3	0.002292	0.001599	1.434	0.15293
x4	4.022544	0.021194	189.799	< 2e-16 ***
x5	-0.630767	0.254732	-2.476	0.01395 *
x6	3.775812	0.065497	57.648	< 2e-16 ***
x10	0.935771	0.452770	2.067	0.03979 *

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
## Residual standard error: 0.5987 on 248 degrees of freedom
## Multiple R-squared: 0.9996, Adjusted R-squared: 0.9996
## F-statistic: 8.838e+04 on 7 and 248 DF, p-value: < 2.2e-16

model.5 <- lm(y~x1+x2+x4+x5+x6+x10, data=mcdonalds)

summary(model.5)

##
## Call:
## lm(formula = y ~ x1 + x2 + x4 + x5 + x6 + x10, data = mcdonalds)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.54519 -0.28767  0.03628  0.30153  2.51553
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.116425   0.080822  -1.441  0.15098
## x1           9.112407   0.048987 186.018 < 2e-16 ***
## x2          -0.015342   0.004889  -3.138  0.00190 **
## x4           4.015530   0.020665 194.315 < 2e-16 ***
## x5          -0.642730   0.255134  -2.519  0.01239 *
## x6           3.822813   0.056822  67.277 < 2e-16 ***
## x10          1.166167   0.424186   2.749  0.00641 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.6 on 249 degrees of freedom
## Multiple R-squared: 0.9996, Adjusted R-squared: 0.9996
## F-statistic: 1.027e+05 on 6 and 249 DF, p-value: < 2.2e-16
```

Final Model using Backward Selection:

$$Y = -0.116 + 9.112X_1 - 0.015X_2 + 4.015X_4 - 0.642X_5 + 3.822X_6 + 1.166X_{10}$$

Coefficient of Determination

(Assessing the fit of the model) 99.96 % of the variability in the calories is accounted for in this multiple linear regression model.

Prediction:

For a sample that has a total fat (x1) of 4, a cholesterol (x2) of 6, a carbohydrate (x4) of 8.2, a dietary fiber (x5) of 0.4, a protein (x6) of 3.4, and a iron (x10) of 0.54, we predict the calories (y) to be 82.0.

Estimation:

For samples that have a total fat (x1) of 4, a cholesterol (x2) of 6, a carbohydrate (x4) of 8.2, a dietary fiber (x5) of 0.4, a protein (x6) of 3.4, and a iron (x10) of 0.54, we predict the calories (y) to be 82.0.

Interpretation of Partial Slopes (B-weights or Coefficients):

For a fixed total fat (x1), a fixed cholesterol (x2), a fixed carbohydrates (x4), a fixed dietary fiber (x5), and a fixed protein (x6), as the iron increases by 1, the calories increases by 1.166.

For a fixed total fat (x1), a fixed cholesterol (x2), a fixed carbohydrates (x4), a fixed dietary fiber (x5), and a fixed iron (x10), as the protein increases by 1, the calories increases by 3.822.

For a fixed total fat (x1), a fixed cholesterol (x2), a fixed carbohydrates (x4), a fixed protein (x6), and a fixed iron (x10), as the dietary fiber increases by 1, the calories decreases by 0.642.

For a fixed total fat (x1), a fixed cholesterol (x2), a fixed dietary fiber (x5), a fixed protein (x6), and a fixed iron (x10), as the carbohydrates increases by 1, the calories increases by 4.015.

For a fixed total fat (x1), a fixed carbohydrates (x4), a fixed dietary fiber (x5), a fixed protein (x6), and a fixed iron (x10), as the cholesterol increases by 1, the calories decreases by 0.015.

For a fixed cholesterol (x2), a fixed carbohydrates (x4), a fixed dietary fiber (x5), a fixed protein (x6) and a fixed iron (x10), as the total fat increases by 1, the calories increases by 9.112.

Confidence Intervals for the Coefficients

```
confint(model.5, level=0.95)
```

##	2.5 %	97.5 %
## (Intercept)	-0.27560643	0.042756083
## x1	9.01592622	9.208888169
## x2	-0.02497068	-0.005713673
## x4	3.97482982	4.056231109
## x5	-1.14522649	-0.140233825
## x6	3.71090073	3.934725261
## x10	0.33071590	2.001617540

We are 95% confident that for a fixed total fat (x1), a fixed cholesterol (x2), a fixed carbohydrates (x4), a fixed dietary fiber (x5), and a fixed protein (x6), as the iron increases by 1, the calories increases between 0.331 and 2.001.

We are 95% confident that for a fixed total fat (x1), a fixed cholesterol (x2), a fixed carbohydrates (x4), a fixed dietary fiber (x5), and a fixed iron (x10), as the protein increases by 1, the calories increases between 3.711 and 3.935.

We are 95% confident that for a fixed total fat (x1), a fixed cholesterol (x2), a fixed carbohydrates (x4), a fixed protein (x6), and a fixed iron (x10), as the dietary fiber increases by 1, the calories decreases between -1.145 and -0.140.

We are 95% confident that for a fixed total fat (x1), a fixed cholesterol (x2), a fixed dietary fiber (x5), a fixed protein (x6), and a fixed iron (x10), as the carbohydrates increases by 1, the calories increases between 3.975 and 4.056.

We are 95% confident that for a fixed total fat (x1), a fixed carbohydrates (x4), a fixed dietary fiber (x5), a fixed protein (x6), and a fixed iron (x10), as the cholesterol increases by 1, the calories decreases between -0.025 and -0.006.

We are 95% confident that for a fixed cholesterol (x2), a fixed carbohydrates (x4), a fixed dietary fiber (x5), a fixed protein (x6) and a fixed iron (x10), as the total fat increases by 1, the calories increases between 9.015 and 9.209.

Residual Analysis

This is checking the assumptions that need to be satisfied before it is appropriate to perform inference from a multiple regression model.


1. The random errors are independent of each other.
2. The random errors are normally distributed



3. The random errors have constant variance (homoscedasticity)


(Note: There is a lot more to residual analysis than just these things, but to keep things focused on the inferential aspects of linear regression we will just do this quick check.)

```
bptest(y~x1+x2+x4+x5+x6+x10, varformula = ~ fitted.values(model.5), studentize=TRUE, data=mcdonalds)
```

```
##  
## studentized Breusch-Pagan test  
##  
## data: y ~ x1 + x2 + x4 + x5 + x6 + x10  
## BP = 21.68, df = 1, p-value = 3.222e-06  
#oldpar <- par(oma=c(0,0,3,0), mfrow=c(2,2))  
plot(model.5)
```







```
#par(oldpar)
```

Ho: The variability of the residuals is constant (homoscedasticity)

Ha: The variability of the residuals is not constant (heteroscedasticity)


Test statistic: $\chi^2_{BP} = 21.68$

P-value: 3.222e-06

Conclusion (at the .05 level): Reject Ho in favor of Ha. There is sufficient evidence to conclude that the variability of the residuals is not constant (heteroscedasticity).

```
# Examination of the distribution of the residuals
```

```
boxplot(model.5$residuals, col="lightblue", horizontal = TRUE)
```



```
shapiro.test(model.5$residuals)
```

```
##
##  Shapiro-Wilk normality test
##
## data:  model.5$residuals
## W = 0.94993, p-value = 1.066e-07
```

Ho: The residuals follow a normal distribution

Ha: The residuals do not follow a normal distribution

Test statistic: $W^* = 0.94993$

P-value: 1.066e-07

Conclusion (at the .05 level): Reject Ho in favor of Ha. There is sufficient evidence to conclude that the residuals do not follow a normal distribution.

In multiple linear regression, there are several assumption: the residuals (errors) should be normally distributed and should have constant variance. However, in our case, residuals do not follow a normal distribution and have non-constant variance. This suggests that our model may not be a good fit for the Macdonald's menu data. We could explore alternative regression techniques or include interaction terms in the model to better capture the relationships in the data.